

---

# Transaction Management

F28DM Database Management Systems

Matthew P Aylett

[m.aylett@hw.ac.uk](mailto:m.aylett@hw.ac.uk)

Materials from: Alasdair G J Gray

# What is a transaction?

- Series of operations to perform a task
- Logical unit of work



[http://www.strategicdealslawblog.com/files/2012/09/dramstime\\_m\\_8745845.jpg](http://www.strategicdealslawblog.com/files/2012/09/dramstime_m_8745845.jpg)

# Example: Money Transfer

Move £100 between two accounts

What would happen if only the first query executed?

R/W 1. UPDATE Account SET  
balance = balance - 100  
WHERE accountNo = 123;

R/W 2. UPDATE Account SET  
balance = balance + 100  
WHERE accountNo = 124;

# Example: Web Purchase

1. Create order
2. Add customer details to order
3. Place items from basket into order
4. Take payment



# Why use transactions?

- Ensure data integrity
  - Require multiple updates to appear as one
  - Ensure integrity constraints, e.g. referential integrity (FK)
- Support concurrent access
  - Allow multiple users!
  - More throughput
- Recovery, e.g. system crash
  - Avoid data loss



<https://d3glfbbr3jeumb.cloudfront.net/assets/features/concurrent-connections-208bcc5d5db456d66914e808c42a4c05.png>

# User Interactions

- 
- A close-up photograph of a small, light-colored snail with a yellowish-brown shell moving slowly across a black computer keyboard. The snail is positioned in the center of the frame, with its head and antennae extended. The keyboard keys are visible, including 'W', 'E', 'R', 'Z', 'U', 'I', 'S', 'D', 'H', 'J', 'K', 'X', 'C', 'V', 'B', 'N', and 'M'. The background is blurred, showing more of the keyboard and a dark surface. The overall tone is somewhat somber due to the blue tint.
- <https://articles-images.sftcdn.net/wp-content/uploads/sites/3/2017/03/dgddgd-1024x576.jpg>

<https://articles-images.sftcdn.net/wp-content/uploads/sites/3/2017/03/dgddgd-1024x576.jpg>

*Causes problems for transaction processing!*

# Problem 1: Lost Update

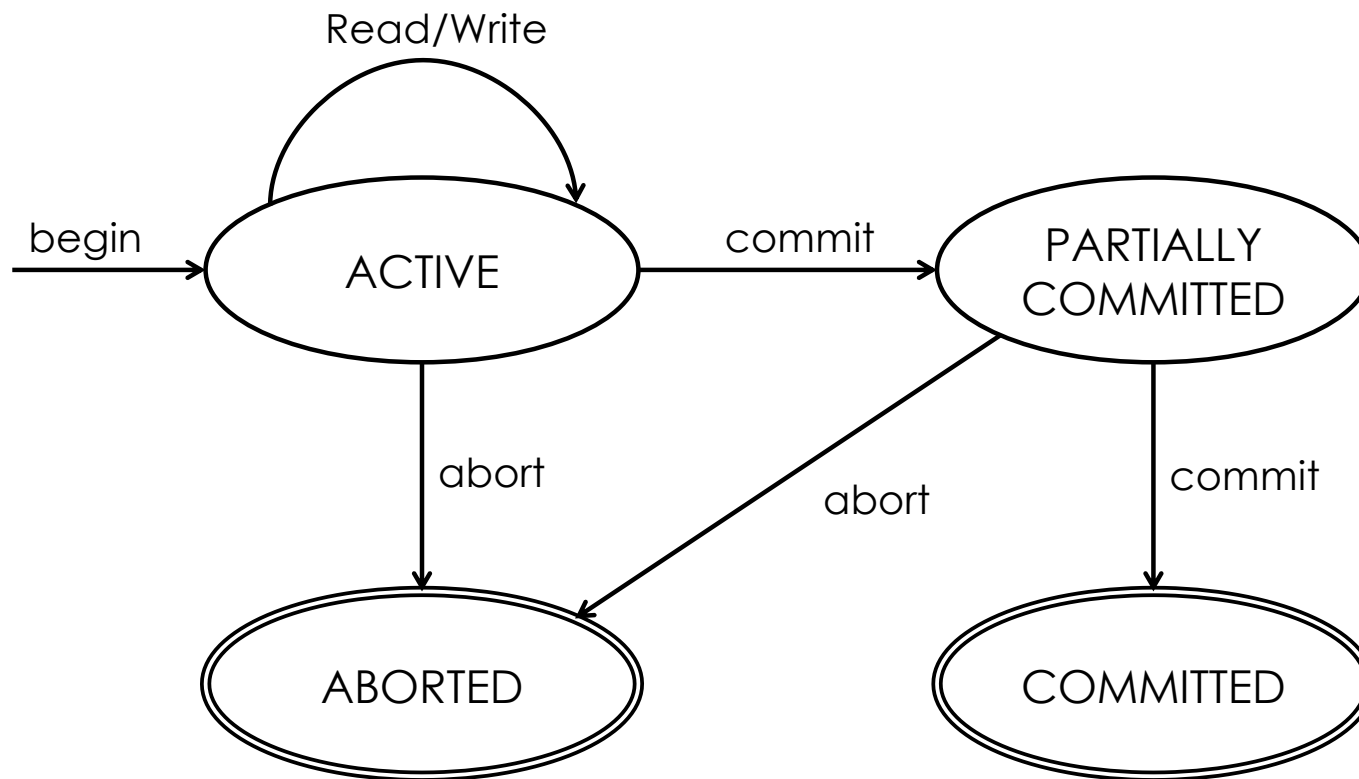
Transaction A	Transaction B	Value V
Get V		5
	Get V	5
Add 10		5
	Add 15	5
Put V		
	Put V	

# Transaction Properties: ACID

- **Atomicity:** All actions complete or none
  - All or nothing
- **Consistency:** Database finishes in a consistent state, i.e. no integrity constraints are violated
  - Only valid data is saved
- **Isolation:** No interference between concurrent transactions
  - Transactions do not affect each other
- **Durability:** Changes are permanent
  - Written data will not be lost

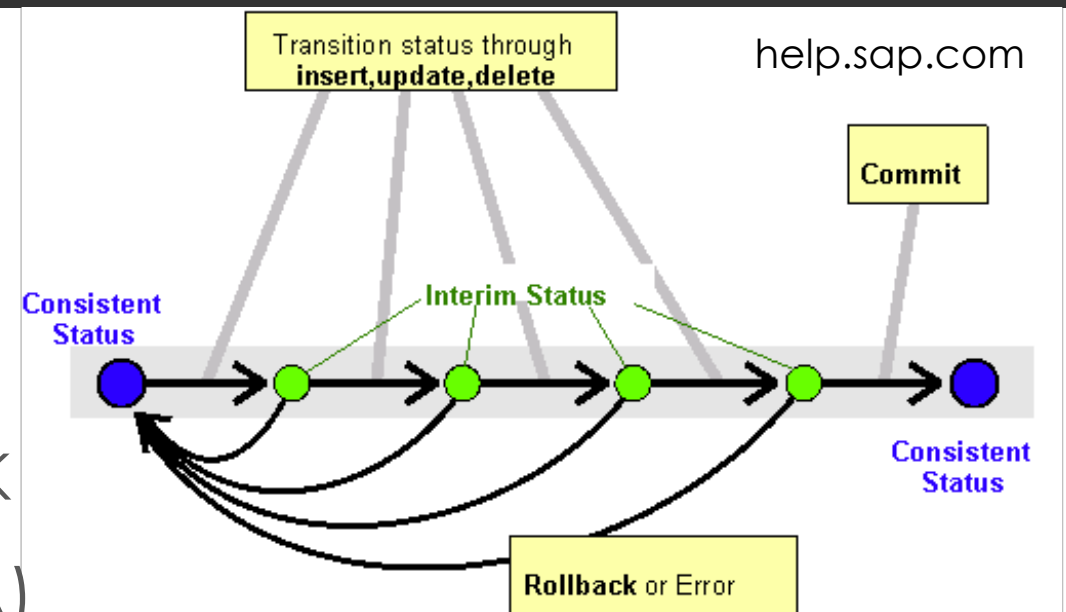


# Transaction Management State Diagram



# Rollback

- ❑ Transaction is aborted:
  - ❑ Failed operation
  - ❑ Error
  - ❑ Failed integrity constraint
  - ❑ Timeout
  - ❑ SQL command: ROLLBACK
- ❑ All operations undone (A)
  - ❑ Appear like the transaction never took place
- ❑ Must not interfere with other transactions (I)



# SQL: Transaction Control Language

## Control statements

▣ <http://dev.mysql.com/doc/refman/5.6/en/commit.html>

▣ **BEGIN:** Start a transaction

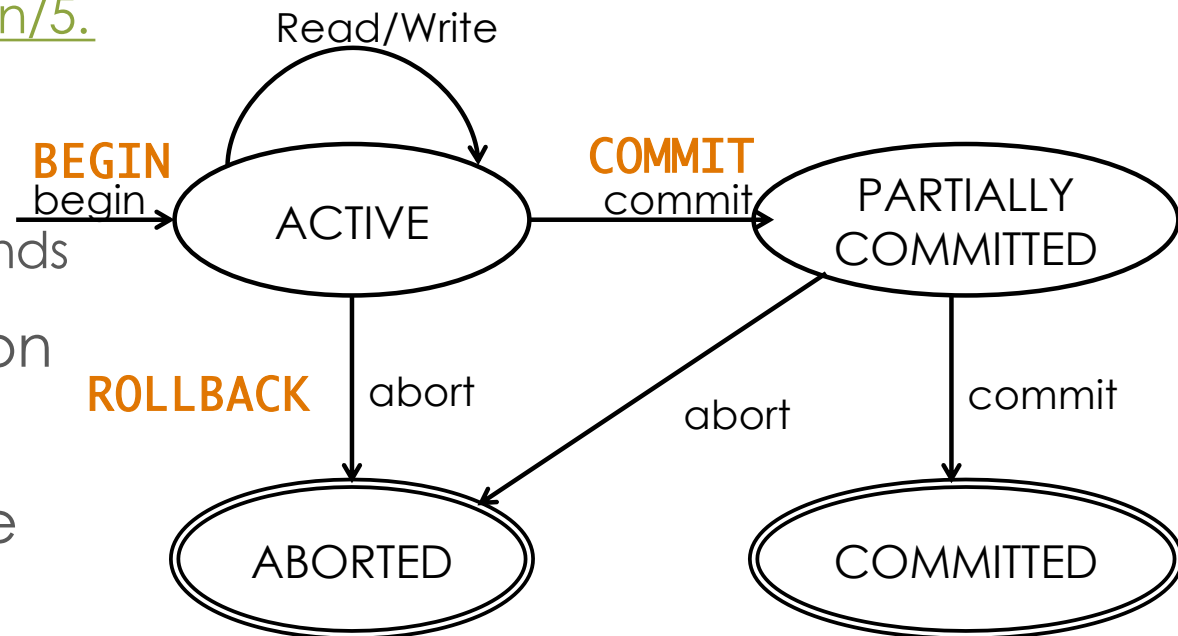
▣ Issue prior to running SQL commands

▣ **COMMIT:** Complete a transaction

▣ Issue once all SQL commands run

▣ **ROLLBACK:** Undo the work since last commit

▣ Issue to return to previous state



# Schedule (Definitions)

- **Schedule:** Set of read and write operations from one or more transactions
  
- **Serial schedule:** Set of read and write operations from a one transaction take place before any operations of another transaction.  
*i.e. All operations of T1, then all operations of T2*
  - No interleaving of operations
  - Low throughput
  - No ACID violations

# Transaction Processing

*Aim to interleave operations from multiple transactions*

- ▣ Maximise database throughput
- ▣ **Serializable schedule:** Interleaving of operations from more than one transaction so that it appears as if one transaction takes place before another.
  - ▣ High throughput
  - ▣ Potential for ACID violations
  - ▣ Equivalent serial schedule, but less efficient
  - ▣ For  $n$  transactions there are  $n!$  possible serial schedules
    - ▣ Unfeasible to check them all

# Schedules

## Serial Schedule

Transaction A	Transaction B
Get V1	
Add 5	
Put V1	
Get V2	
Add 5	
Put V2	
	Get V1
	Get V2


## Serialisable Schedule

Transaction A	Transaction B
Get V1	
Add 5	
Put V1	
	Get V1
Get V2	
Add 5	
Put V2	
	Get V2

# Serialization Graph

## ■ Serializability Theorem:

*A schedule  $S$  is serializable if and only if its corresponding serializability graph  $SG(S)$  is acyclic*




Algorithm to construct a serializability graph for a set of transactions  $\mathcal{T}$

```
for  $T \in \mathcal{T}$  do  
    create node  
end for  
for each operation do  
    if ( $R \in T'$  follows  $W \in T$ ) or  
    ( $W \in T'$  follows  $\{R|W\} \in T$ ) then  
        create edge from  $T$  to  $T'$   
    end if  
end for
```

# Drawing Serialization Graphs


## Serializable Schedule

Transaction A	Transaction B
Read V	
Write V	
	Read V
	Write V



## Non-serializable Schedule

Transaction A	Transaction B
Read V	
	Read V
Write V	
	Write V






# Drawing

Read X Follows Write Y then  $Y \rightarrow X$   
 Write X follows Read Y then  $Y \rightarrow X$   
 Write X Follows Write Y then  $Y \rightarrow X$


## Serializable Schedule

Transaction A	Transaction B
Read V	
Write V	
	Read V
	Write V



## Non-serializable Schedule

Transaction A	Transaction B
Read V	
	Read V
Write V	
	Write V



# Homework

□ Draw serialisation graph for problem 3

Transaction A	Transaction B	Value V1	Value V2
Get V1		5	15
Add 5		5	15
Put V1		10	15
	Get V1	10	15
	Get V2	10	15
Get V2		10	15
Add 5		10	15
Put V2		10	20

# SQL Transaction

- Transaction to move money between accounts
- Don't want to “lose” money
  - Atomic: All should happen or none – money lost in ether
  - Consistency: IC to ensure account balance > 0
  - Isolation: Should not interfere with other transactions – money lost to ether
  - Durable: Once committed change not be lost – system crash

1.	START TRANSACTION;
2.	UPDATE Account SET balance = balance - 100 WHERE accountNo = 86036243;
3.	UPDATE Account SET balance = balance + 100 WHERE accountNo = 78612361;
4.	[COMMIT   ROLLBACK];

# SQL Transaction with Variable

■ Transaction all of Gareth Scarth's money to Zoe Kender

■ Use variables to capture account number and balance

1.	START TRANSACTION;
2.	SELECT @a1:= accountNumber, @b1:=balance FROM Account WHERE firstnames = 'Gareth' AND lastname = 'Scarth'
3.	UPDATE Account SET balance = 0 WHERE accountNo = @a1;
4.	UPDATE Account SET balance = balance + @b1 WHERE firstnames = 'Zoe' AND lastname = 'Kender';
5.	[COMMIT   ROLLBACK];

# SQL Isolation Levels for Transactions

- **Serializable** (highest): Appears that a transaction run entirely before or after others.
- **Read repeatable** (InnoDB default): Repeated read gets same data or same data with new inserts.
- **Read committed**: Repeated reads get new values from committed transactions.
- **Read uncommitted** (“dirty read”): Reads can see values from other concurrent transactions

# Summary

- **Transactions:** group operations into unit of work
- Provide **ACID guarantees**
  - **Atomicity:** All or nothing
  - **Consistency:** Valid against schema and constraints
  - **Isolation:** Does not interfere with other transactions
  - **Durable:** Once committed actions are not lost
- Enable concurrent access
- Interleave operations into **Serialisable Schedule**

---

# Concurrency Control

F28DM Database Management Systems

Matthew P Aylett

[m.aylett@hw.ac.uk](mailto:m.aylett@hw.ac.uk)

Materials from: Alasdair G J Gray

# Revision Quiz

- What is a transaction?
- What properties does the transaction manager guarantee?
- What is a serializable schedule?



# Example: Money Transfer

Move £100 between two accounts

What would happen if only the first query executed?

R/W 1. UPDATE Account SET  
balance = balance - 100  
WHERE accountNo = 123;

R/W 2. UPDATE Account SET  
balance = balance + 100  
WHERE accountNo = 124;

# Transaction Properties: ACID

- **Atomicity:** All actions complete or none
  - All or nothing
- **Consistency:** Database finishes in a consistent state, i.e. no integrity constraints are violated
  - Only valid data is saved
- **Isolation:** No interference between concurrent transactions
  - Transactions do not affect each other
- **Durability:** Changes are permanent
  - Written data will not be lost

# Schedules

## Serial Schedule

Transaction A	Transaction B
Get V1	
Add 5	
Put V1	
Get V2	
Add 5	
Put V2	
	Get V1
	Get V2

## Serialisable Schedule

Transaction A	Transaction B
Get V1	
Add 5	
Put V1	
	Get V1
Get V2	
Add 5	
Put V2	
	Get V2

# Concurrency Control

Running many transactions in parallel

- There are many strategies because how efficient it is depends on the way data is accessed and especially written
- That depends on what your database is for
- Different strategies have different advantages and disadvantages

# Concurrency Control

Running many transactions in parallel

- **Single-Version:** maintains single copy of data
  - Locking (pessimistic):
    - Assumes conflicts
    - Locks prevent interactions
  - Timestamp (optimistic):
    - Deal with conflicts when they happen
    - Repeat work
- **Multi-Version:** maintains multiple copies of data

# Locks

- Shared Lock (s-lock)  
aka read-lock:
  - Required for reading data
  - Many transactions can hold simultaneously

- Exclusive Lock (X-lock)  
aka write-lock:
  - Required for writing data
  - Only one transaction
  - No shared locks can exist for other transactions

## Lock Granting Matrix


Lock Held	Shared requested	Exclusive requested
Shared	Granted	Rejected
Exclusive	Rejected	Rejected

# Lock Granularity


Table




Disk Block



Row



Cell




# Blocking & Deadlock

■ **Blocking:** waiting to acquire lock held by another transaction

■ **Deadlock:** Cycle of transactions waiting for locks


T <sub>A</sub>	T <sub>B</sub>
Get V1	
Put V1	
	Get V2
	Get V1
Get V2	
Put V2	





# Handling Deadlock

- **Time-out:** Abort blocked transaction
  - Simple to implement, pick random transaction
  - Aggressive: can undo more than required
- **Detection:** Identify cycle
  - Draw wait-for graph (serialisation graph)
    - Requires processing
  - Abort selected transactions
    - Don't undo unnecessarily
- **Two phase locking** protocol



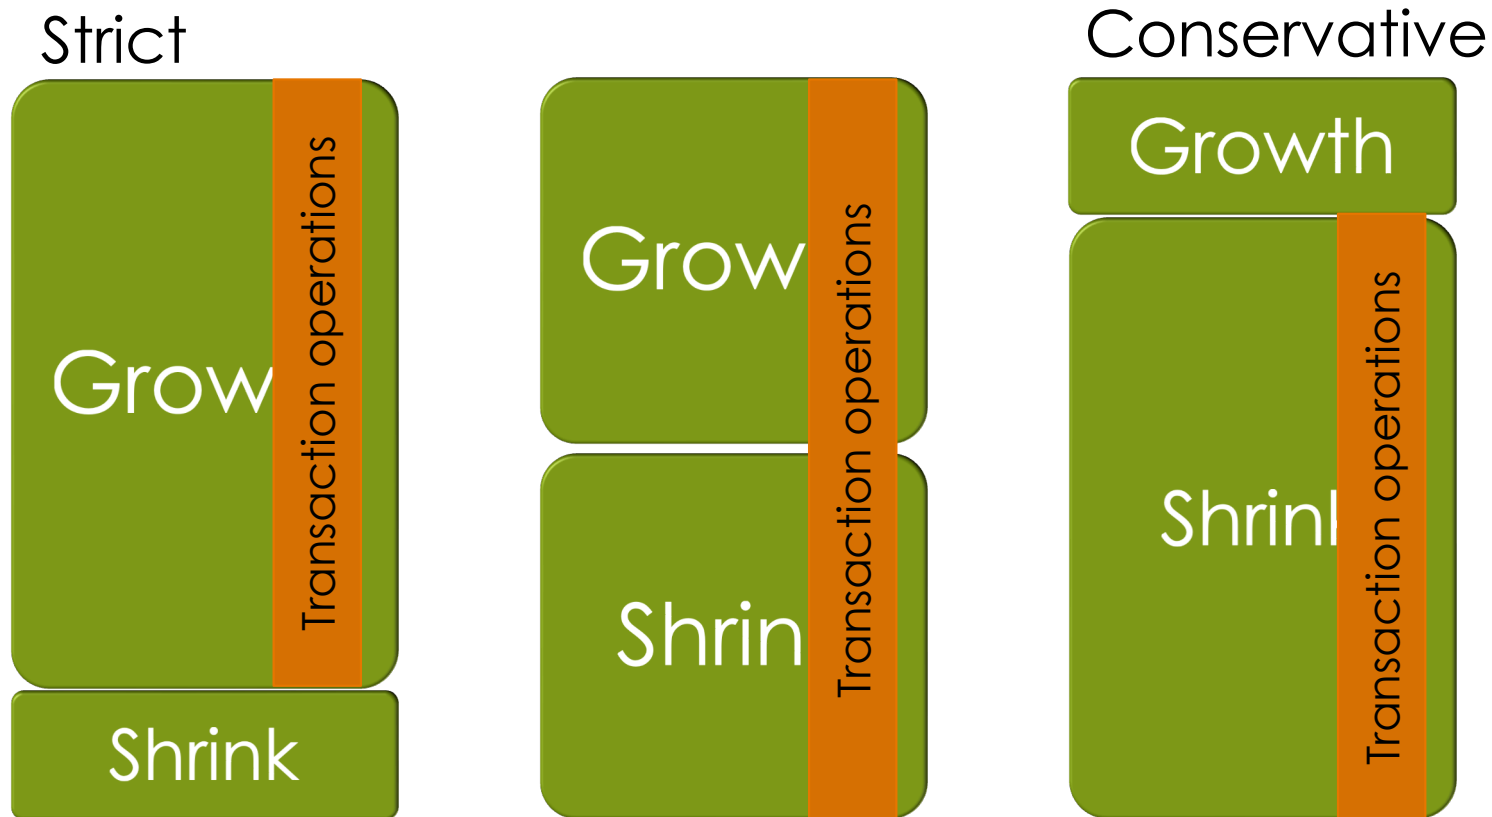
<http://pages.cs.wisc.edu/~bart/537/lecturenotes/figures/s12.crash.gif>

# Two-Phase Locking (2PL)



- **Growth phase:** transaction acquires locks as needed
- **Shrink phase:** transaction releases locks
  - Can no longer acquire locks
- **2PL:** Locks acquired as needed, locks released when no longer required and no more locks will be needed
  - Problems occur if a lock released and transaction aborted

# Two-Phase Locking Summary



# 2PL: Example



T <sub>A</sub>	T <sub>B</sub>	2PL
BEGIN		
Get V1		T <sub>A</sub> : Acquire S-lock on V1
Put V1		T <sub>A</sub> : Acquire X-lock on V1 Lock not released as require more locks
	BEGIN	
	<del>Get V1</del>	T <sub>B</sub> : blocked acquiring S-lock on V1, transaction paused
Get V2		T <sub>A</sub> : Acquire S-lock on V2
Put V2		T <sub>A</sub> : Acquire X-lock on V2.
Get V1		T <sub>A</sub> : Lock on V2 released, already have lock on V1
COMMIT		T <sub>A</sub> : Release lock on V1
	GET V1	T <sub>B</sub> : Unpaused. Acquire S-lock on V1
	...	

# 2PL: Read Uncommitted (dirty read)



T <sub>A</sub>	T <sub>B</sub>	2PL
BEGIN		
Get V1		T <sub>A</sub> : Acquire S-lock on V1
Put V1		T <sub>A</sub> : Acquire X-lock on V1 Lock not released as require more locks
	BEGIN	
	<del>Get V1</del>	T <sub>B</sub> : blocked acquiring S-lock on V1, transaction paused
Get V2		T <sub>A</sub> : Acquire S-lock on V2
Put V2		T <sub>A</sub> : Acquire X-lock on V2, release lock on V1
	GET V1	T <sub>B</sub> : Unpaused. Acquire S-lock on V1
ABORT		T <sub>A</sub> : Work undone, release lock on V2
	...	

T<sub>B</sub> has read the updated value of T<sub>A</sub>

# Strict Two-Phase Locking



- ▣ Locks acquired as needed
- ▣ All locks kept until the end of the transaction
  - ▣ Released in one go
  - ▣ Eliminates 'dirty read' problem of 2PL
  - ▣ Reduces throughput

# Strict Two-Phase Locking



T <sub>A</sub>	T <sub>B</sub>	2PL
BEGIN		
Get V1		T <sub>A</sub> : Acquire S-lock on V1
Put V1		T <sub>A</sub> : Acquire X-lock on V1
	BEGIN	
	<del>Get V1</del>	T <sub>B</sub> : blocked acquiring S-lock on V1, transaction paused
Get V2		T <sub>A</sub> : Acquire S-lock on V2
Put V2		T <sub>A</sub> : Acquire X-lock on V2.
Get V1		T <sub>A</sub> : Already have lock on V1
COMMIT		T <sub>A</sub> : Release locks on V1 & V2
	GET V1	T <sub>B</sub> : Unpaused. Acquire S-lock on V1
	...	

# Strict 2PL: No dirty read



T <sub>A</sub>	T <sub>B</sub>	2PL
BEGIN		
Get V1		T <sub>A</sub> : Acquire S-lock on V1
Put V1		T <sub>A</sub> : Acquire X-lock on V1 Lock not released as require more locks
	BEGIN	
	<del>Get V1</del>	T <sub>B</sub> : blocked acquiring S-lock on V1, transaction paused
Get V2		T <sub>A</sub> : Acquire S-lock on V2
Put V2		T <sub>A</sub> : Acquire X-lock on V2, <b>lock on V1 is not released</b>
ABORT		T <sub>A</sub> : Work undone, <b>release locks on V1 and V2</b>
	GET V1	T <sub>B</sub> : Unpaused. Acquire S-lock on V1
	...	

T<sub>B</sub> cannot proceed until T<sub>A</sub> completes



# Conservative Two-Phase Locking



- All locks acquired at start
- Locks released once no longer needed
  - Locks released before end of transaction
  - Susceptible to dirty read problem
  - Prevents deadlocks
  - Better throughput than strict

# Example: Conservative 2PL



T <sub>A</sub>	T <sub>B</sub>	2PL
BEGIN		T <sub>A</sub> : Acquire X-locks on V1 & V2
Get V1		
Put V1		
	BEGIN	T <sub>B</sub> : blocked acquiring X-lock on V1, transaction paused
Get V2		
Put V2		
Get V1		T <sub>A</sub> : Release lock on V2
	GET V1	T <sub>B</sub> : Unpaused. Having acquired S-lock on V1
COMMIT		T <sub>A</sub> : Release lock V2
	...	

# Conservative 2PL: Dirty read

Growth

Shrink

Transaction operations

T <sub>A</sub>	T <sub>B</sub>	2PL
BEGIN		T <sub>A</sub> : Acquire X-locks on V1 & V2
Get V1		
Put V1		
	BEGIN	T <sub>B</sub> : blocked acquiring X-lock on V1, transaction paused
Get V2		
Put V2		Release lock on V1
	GET V1	T <sub>B</sub> : Unpaused. Acquires required locks ( <i>T<sub>B</sub> does not use V2</i> )
ABORT		T <sub>A</sub> : Work undone, release lock on V2
	...	

T<sub>B</sub> has read the updated value of T<sub>A</sub>

# Homework: MySQL Locks

Perform operations in two different terminal windows both connected to MySQL. Terminal 2 cannot proceed until Terminal 1 commits.

Operation	Terminal 1	Terminal 2
1	START TRANSACTION;	
2	UPDATE account SET balance = 100 WHERE accountNo = 23525;	
3		UPDATE account SET balance = balance + 10 WHERE accountNo = 23525;
4	UPDATE account SET balance = 100 WHERE accountNo = 23526;	
5	COMMIT;	

# Timestamp Protocol

- Each transaction assigned timestamp  $t_T$
- Each data item has
  - Last read timestamp  $t_r$
  - Last write timestamp  $t_w$
- Operation permitted iff
$$t_{r|w} < t_T$$
- Otherwise
  - Abort transaction  $t_T$
  - Restart transaction

Exercise: Draw sequence of events for the three problems



# Timestamp Protocol

Vs timestamp

Transaction 28	Transaction 32
Read V	
	Read V
Write V	
	Write V

Read

26

Write

26

# Multi-version Concurrency (MVCC)

- Multiple simultaneous transactions
- Each transaction sees isolated snapshot
- Changes only seen across transactions on commit

# MVCC in Action

xmin	xmax	name	notes
100	0	Alice	Great at programming
101	0	Bob	Always talking to Alice
102	0	Eve	Listens to everyone's conversations

## ■ TXID 103: update Bob

xmin	xmax	name	notes
100	0	Alice	Great at programming
101	<b>103</b>	Bob	Always talking to Alice
102	0	Eve	Listens to everyone's conversations
103	0	Bob	Working very hard

## ■ TXID 102: long running read over entire table

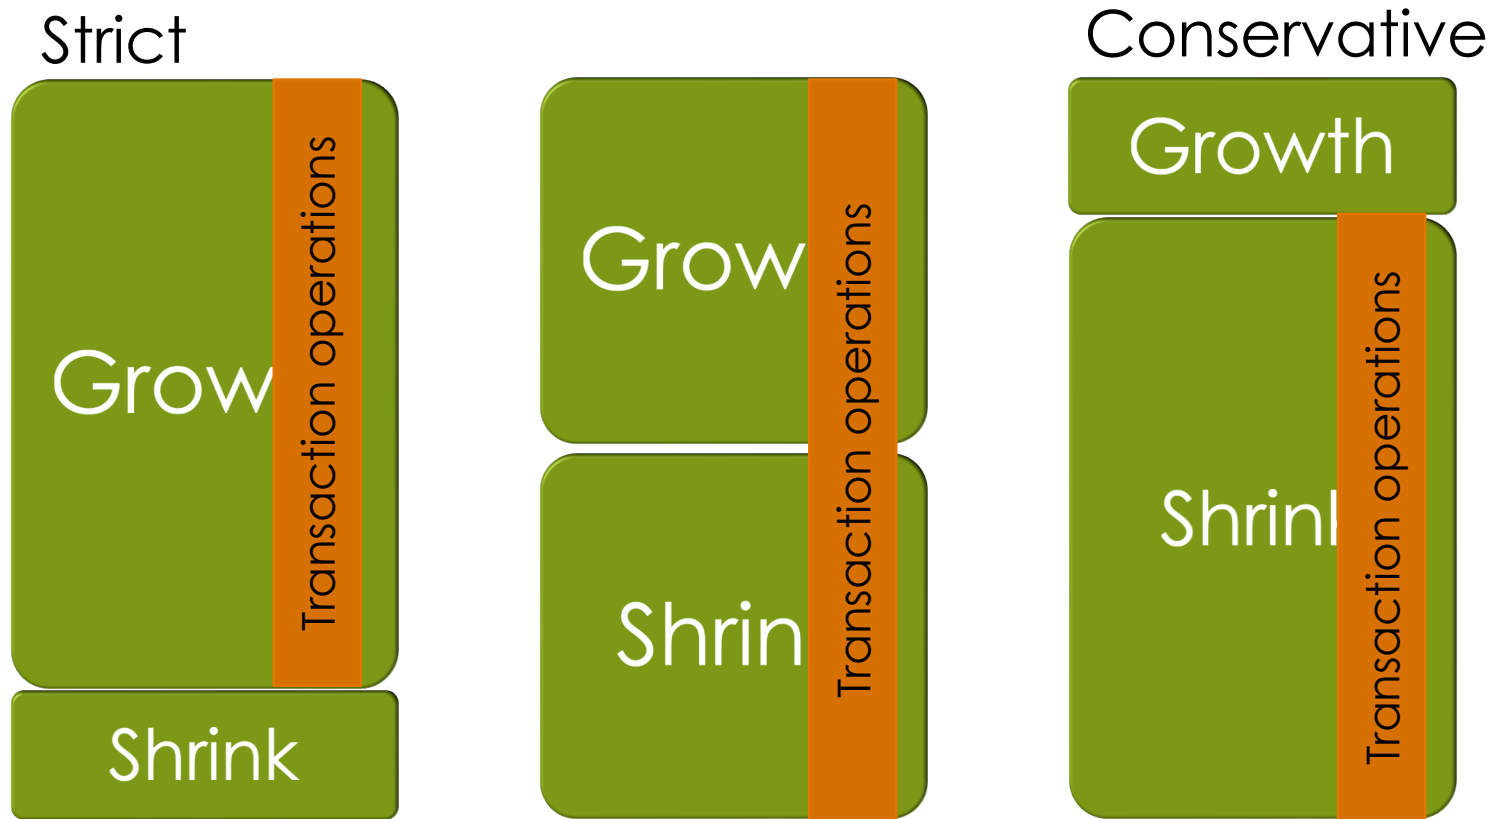
- Reads value 'Always talking to Alice'
- Won't see value about Eve (only see values strictly less than TID)



# Summary

- **Concurrency control:** enables multiple users to interact with the database
- **2 phase commit:** uses **locks** (s-lock/x-lock)
  - Classical: dirty read problem
  - Strict: reduces throughput
  - Conservative: dirty read problem
- **Timestamp concurrency control:** can only operate on old values
- **Multi-version concurrency control:** retains multiple versions of the data and when they are valid

# Two-Phase Locking Summary



# References

Connolly, T., & Begg, C. (2005). *Database Systems: A Practical Approach to Design, Implementation, and Management* (4th ed.). Addison Wesley. Chapter 20

Ward, P. (2008). *Database Management Systems* (2nd ed.). Middlesex University Press. Chapter 9