



C Programming

Casting

Adam Sampson (based on material by Greg Michaelson)

School of Mathematical and Computer Sciences

Heriot-Watt University

C lectures

- Compiling code, program layout, printing/reading data, expressions, arithmetic, memory addresses, control flow, precedence
- Functions, pointers, file IO, arrays
- **Memory allocation, casting, masking, shifting**
- Strings, structures, dynamic space allocation, field access
- Recursive structures, 2D arrays, union types

Recap: `sizeof` and arrays

- `sizeof(type)` tells you how much memory a variable of that type would occupy
- `sizeof(char)` is **always** 1
 - C counts sizes in chars – usually 8-bit bytes
- Arrays are contiguous in memory
- Declaring array:
`type name[length];`
allocates size `length * sizeof(type)`

Low-level programming in C

- We often need to care about how exactly things are laid out in memory in C...
- e.g. controlling hardware devices, using hardware registers that are mapped into memory at specific addresses
- e.g. writing code to do custom memory allocation

Declarations and space

- Amount of space for a given type depends on:
 - Type – e.g. `char` vs `float`
 - Platform – e.g. 64-bit x86, 32-bit ARM
 - Compiler and version – for structure types (primitive types are normally standardised)
- `sizeof(type)`'s return type is **`size_t`**
 - An unsigned integer type big enough to hold any size – to print with `printf`, use `%zd`

typesize.c

Size of standard types

Sizes of types

```
#include <stdio.h>
int main(int argc, char *argv[])
{
    printf("char:      %zd\n", sizeof(char));
    printf("short:     %zd\n", sizeof(short));
    printf("int:        %zd\n", sizeof(int));
    printf("long:       %zd\n", sizeof(long));
    printf("float:      %zd\n", sizeof(float));
    printf("double:    %zd\n", sizeof(double));
    printf("char *:    %zd\n", sizeof(char *));
    return 0;
}
```

Typical type sizes

64-bit x86 PC

```
$ ./typesize
```

```
char:    1
```

```
short:   2
```

```
int:     4
```

```
long:    8
```

```
float:   4
```

```
double:  8
```

```
char *:  8
```

32-bit ARM RPi

```
$ ./typesize
```

```
char:    1
```

```
short:   2
```

```
int:     4
```

```
long:    4
```

```
float:   4
```

```
double:  8
```

```
char *:  4
```

64-bit RISC-V

```
$ ./typesize
```

```
char:    1
```

```
short:   2
```

```
int:     4
```

```
long:    8
```

```
float:   4
```

```
double:  8
```

```
char *:  8
```


Explicitly-sized types

- What if you definitely want a 64-bit type, regardless of the platform?
- C99: `<stdint.h>` header defines types like:
 - `int64_t` – signed 64-bit integer
 - `uint32_t` – unsigned 32-bit integer
- Also macros for `printf` formats for these – see the documentation

Recap: bits, bytes, hexadecimal


- A hexadecimal constant in C:
 $0xh_1h_2 \dots h_N$
where $h_i == 0..9 A..F$ (or a..f)
- Each hex digit is 4 bits, so an 8-bit byte is 2 hex digits
- In $0xAB$, A is bits 7-4, B is bits 3-0
- e.g. $0xFF == 1111\ 1111 ==$ all bits 1
 $0x00 == 0000\ 0000 ==$ all bits 0
 $0x65 == 0110\ 0101 ==$ ASCII 'e'

Big- and little-endian machines


- When a value occupies more than one byte, what order are the bytes stored in? Depends on the machine...

e.g. `int a; // 32 bit type`
`a = 0x01234567;`

- Little-endian** architecture – least significant byte stored first



- Big-endian** architecture – most significant byte stored first




Big- and little-endian machines

- When a value occupies more than one byte, what order are the bytes stored in? Depends on the machine...

e.g. `int a; // 32 bit type`
`a = 0x01234567;`

- Little-endian** architecture – least significant byte stored first



Nearly all modern architectures are little-endian.

(ARM can be switched to operate in either mode, but it's nearly always used in little-endian mode.)

Accessing bytes in C


- In C, we can access the individual bytes that make up a value by **pointer casting**
 - e.g. we can take an `int`, and reinterpret it as an array of 4 `chars`
- Casting is C's explicit **type conversion** mechanism
 - *(newtype) expression* – evaluate *expression*, then convert the result to *newtype*
 - e.g. `(int) sqrt(47.3)`

cast.c

Casting int to char


Casting integers

```
int x = 0x01234567;
```



Casting integers


```
int x = 0x01234567;  
char *c1 = (char *) &x;
```



Casting integers

```
int x = 0x01234567;  
char *c1 = (char *) &x;  
printf("%x\n", *c1);
```


Prints: 67



Casting integers

```
int x = 0x01234567;  
char *c1 = (char *) &x;  
printf("%x %x %x %x\n",  
       c1[0], c1[1], c1[2], c1[3]);
```

Prints: 67 45 23 1



Strict aliasing

- The previous example is legal C, but many similar pieces of code wouldn't be
 - e.g. `float *f = (float *) &a;`
- Standard C's **aliasing** rules say that you aren't allowed to refer to a value by a pointer type that doesn't match the original declaration
- ... **unless** the pointer type is `char *` – which is why the previous example's OK!

Strict aliasing

- Many C programmers are not aware of this rule, and you'll see incorrect “type punning” code in **many** real programs
- In most cases it'll work, but modern compilers are getting more aggressive about taking advantage of aliasing when optimising code...
- ... so compilers often have an option to make it legal, e.g. GCC's `-fno-strict-aliasing`